

603 final code

December 1, 2021

```
[39]: import pyspark
from pyspark.sql import SparkSession
from pyspark.sql.types import FloatType, StructType, StringType, StructField,
↳ IntegerType
from pyspark.ml.regression import LinearRegression
from pyspark.ml.feature import VectorAssembler
import matplotlib.pyplot as plt
from matplotlib.pyplot import figure
import numpy as np
```

```
[40]: spark = SparkSession.builder.getOrCreate()
sc = spark.sparkContext

combined_schema = StructType([
    StructField('STATION', StringType()),
    StructField('NAME', StringType()),
    StructField('LATITUDE', FloatType()),
    StructField('LONGITUDE', FloatType()),
    StructField('ELEVATION', FloatType()),
    StructField('AWND', FloatType()),
    StructField('PRCP', FloatType()),
    StructField('SNOW', FloatType()),
    StructField('SNWD', FloatType()),
    StructField('TAVG', FloatType()),
    StructField('CRASHCOUNT', IntegerType()),
    StructField('YEAR', IntegerType()),
    StructField('MONTH', IntegerType()),
    StructField('DAY', IntegerType())
])
```

```
[41]: combinedDF = spark.read.options(header = 'True').schema(combined_schema).
↳ csv('combined.csv')
combinedDF = combinedDF.drop('STATION', 'NAME', 'ELEVATION', 'LATITUDE',
↳ 'LONGITUDE')
# elevation, latitude and longitude have a large number of missing values
combinedDF.show(30)
```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+

AWND	PRCP	SNOW	SNWD	TAVG	CRASHCOUNT	MONTH	DAY
10.74	0.0	0.0	0.0	50.0	49	4	1
8.28	0.0	0.0	0.0	40.0	35	4	2
10.51	0.0	0.0	0.0	40.0	49	4	3
4.25	0.04	0.0	0.0	37.0	38	4	4
8.05	0.11	0.0	0.0	48.0	60	4	5
5.59	0.0	0.0	0.0	47.0	49	4	6
5.82	0.0	0.0	0.0	50.0	52	4	7
4.92	0.0	0.0	0.0	64.0	64	4	8
6.71	0.0	0.0	0.0	73.0	59	4	9
7.61	0.0	0.0	0.0	76.0	70	4	10
9.62	0.0	0.0	0.0	72.0	55	4	11
10.29	0.67	0.0	0.0	55.0	49	4	12
5.14	0.0	0.0	0.0	54.0	53	4	13
5.59	0.0	0.0	0.0	54.0	58	4	14
8.28	0.0	0.0	0.0	56.0	52	4	15
6.04	0.0	0.0	0.0	60.0	46	4	16
5.82	0.0	0.0	0.0	65.0	62	4	17
9.17	0.04	0.0	0.0	63.0	64	4	18
7.83	0.79	0.0	0.0	66.0	73	4	19
10.29	0.07	0.0	0.0	54.0	47	4	20
4.92	0.0	0.0	0.0	47.0	40	4	21
11.18	0.0	0.0	0.0	46.0	52	4	22
6.26	0.0	0.0	0.0	51.0	46	4	23
6.93	0.0	0.0	0.0	56.0	47	4	24
7.16	0.0	0.0	0.0	58.0	65	4	25
4.25	0.0	0.0	0.0	55.0	55	4	26
4.47	0.0	0.0	0.0	55.0	44	4	27
5.37	0.0	0.0	0.0	57.0	54	4	28
7.16	0.26	0.0	0.0	56.0	66	4	29
7.83	0.22	0.0	0.0	56.0	62	4	30

only showing top 30 rows

```
[42]: features = combinedDF.columns

features.pop(5)

vectorAssembler = VectorAssembler(inputCols = features, outputCol = 'features')

vDF = vectorAssembler.transform(combinedDF)
vDF = vDF.select(['features', 'CRASHCOUNT'])

split = vDF.randomSplit([.8, .2], seed = 12345)
```

```

trainingData = split[0]
testData = split[1]

lr = LinearRegression(featuresCol = 'features', labelCol = 'CRASHCOUNT',
    ↪maxIter = 10, regParam = .3, elasticNetParam = .8)

lrModel = lr.fit(trainingData)

```

```

[43]: #print('Coefficients:', str(lrModel.coefficients))

count = 0

print('Coefficients per feature:')
for item in lrModel.coefficients:
    print(features[count] + ': ', item)
    count += 1
print()

trainingSummary = lrModel.summary
print('RMSE:', trainingSummary.rootMeanSquaredError)
print('R squared:', trainingSummary.r2)

```

```

Coefficients per feature:
AWND: 0.6590887741121741
PRCP: 0.0
SNOW: -1.3372764638638435
SNWD: -0.7240295709285356
TAVG: 0.17799762976077024
MONTH: 0.04447601690300508
DAY: -0.05181953403363659

```

```

RMSE: 15.65479033763715
R squared: 0.06788566821269193

```

```

[44]: lr_predictions = lrModel.transform(testData)
lr_predictions.select('prediction', 'CRASHCOUNT', 'features').show(5, False)

from pyspark.ml.evaluation import RegressionEvaluator
lr_evaluator = RegressionEvaluator(predictionCol = 'prediction', labelCol =
    ↪'CRASHCOUNT', metricName = 'r2')
print('R squared on test data:', lr_evaluator.evaluate(lr_predictions))
testResult = lrModel.evaluate(testData)
print('RMSE on test data:', testResult.rootMeanSquaredError)

```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+

```


prediction	CRASHCOUNT	features
60.86242379345681	55	[0.8899999856948853, 0.029999999329447746, 0.0, 0.0, 49.0, 10.0, 14.0]
65.64618057917546	59	[1.340000033378601, 0.7599999904632568, 0.0, 0.0, 75.0, 8.0, 15.0]
57.05293320973557	67	[1.5700000524520874, 0.0, 0.0, 0.0, 25.0, 1.0, 6.0]
58.9684650252729	57	[1.5700000524520874, 0.0, 0.0, 0.0, 40.0, 12.0, 30.0]
59.86519465279066	39	[1.5700000524520874, 0.07999999821186066, 0.0, 0.0, 43.0, 12.0, 23.0]

only showing top 5 rows

R squared on test data: 0.07517689841308306
RMSE on test data: 16.643762849196168

0.1 Slightly worse

```
[45]: pdDF = lr_predictions.toPandas()
fig, ax = plt.subplots()
fig.set_size_inches(14,14)
plt.scatter(pdDF.prediction, pdDF.CRASHCOUNT)
ax.xaxis.set_ticks(np.arange(0, 170, 10))
ax.yaxis.set_ticks(np.arange(0, 170, 10))
plt.ylabel('Actual')
plt.xlabel('Predicted')
ax.xaxis.label.set_size(30)
ax.yaxis.label.set_size(30)
```



```
[46]: trainingData.describe().show()
```

```
+-----+-----+
|summary|      CRASHCOUNT|
+-----+-----+
|  count|             1975|
|   mean|65.54886075949366|
| stddev|16.21894640101575|
|    min|                 4|
|    max|             109|
+-----+-----+
```

```
[47]: from pyspark.ml.regression import DecisionTreeRegressor

dt = DecisionTreeRegressor(featuresCol='features', labelCol='CRASHCOUNT')
dtModel = dt.fit(trainingData)
dtPredictions = dtModel.transform(testData)
dtEvaluator = RegressionEvaluator(
    labelCol="CRASHCOUNT", predictionCol="prediction", metricName="rmse")
rmse = dtEvaluator.evaluate(dtPredictions)
print("Root Mean Squared Error (RMSE) on test data: ", rmse)
```

Root Mean Squared Error (RMSE) on test data: 16.805152604072795

```
[48]: dtModel.featureImportances
```


```
[48]: SparseVector(7, {0: 0.1706, 1: 0.0803, 2: 0.0329, 3: 0.052, 4: 0.3466, 5: 0.198,
6: 0.1196})
```

```
[49]: features
```

```
[49]: ['AWND', 'PRCP', 'SNOW', 'SNWD', 'TAVG', 'MONTH', 'DAY']
```

```
[50]: fig, ax = plt.subplots()
fig.set_size_inches(14,14)
plt.bar(features, dtModel.featureImportances)
plt.xticks(fontsize=25)
plt.yticks(fontsize=14)
```

```
[50]: (array([0. , 0.05, 0.1 , 0.15, 0.2 , 0.25, 0.3 , 0.35, 0.4 ]),
 [Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, ''),
  Text(0, 0, '')[0]])
```



1 It seems that year is by far the best predictor